

Import Libraries

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder
```

Read Data

<https://www.kaggle.com/datasets/fedesoriano/stroke-prediction-dataset>

```
df = pd.read_csv('dataset-stroke.csv')
```

Preview Dataset

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   id                  5110 non-null  int64  
 1   gender              5110 non-null  object  
 2   age                 5110 non-null  float64 
 3   hypertension        5110 non-null  int64  
 4   heart_disease       5110 non-null  int64  
 5   ever_married        5110 non-null  object
```

```

6  work_type      5110 non-null  object
7  Residence_type 5110 non-null  object
8  avg_glucose_level 5110 non-null float64
9  bmi           4909 non-null float64
10 smoking_status  5110 non-null  object
11 stroke        5110 non-null  int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB

```

```
df.isna().sum()
```

```

id                0
gender            0
age              0
hypertension      0
heart_disease     0
ever_married      0
work_type         0
Residence_type    0
avg_glucose_level 0
bmi              201
smoking_status    0
stroke            0
dtype: int64

```

```
df.head()
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12

Clean Data

```
# Drop out ID column
df = df.drop('id',axis=1)
```

```
# Replace Missing Values in column 'bmi' with the average bmi
df.loc[df['bmi'].isnull(),'bmi'] = df.bmi.mean()
```

```
# Clean gender column
df['gender'].value_counts()
```

```
Female    2994
Male      2115
Other         1
Name: gender, dtype: int64
```

```
df = df.drop(df[df['gender'] == 'Other'].index, axis=0)
```

```
# Convert categorical string into numeric
cols = ['gender', 'ever_married', 'work_type', 'Residence_type', 'smoking_status']

for col in cols:
    print(f"---{col}--- \n{df[col].value_counts()}")
```

```
---gender---
Female    2994
Male      2115
Name: gender, dtype: int64
---ever_married---
Yes      3353
No       1756
Name: ever_married, dtype: int64
---work_type---
Private          2924
Self-employed    819
children         687
Govt_job         657
Never_worked     22
Name: work_type, dtype: int64
---Residence_type---
Urban          2596
```

```
Rural      2513
Name: Residence_type, dtype: int64
---smoking_status---
```

```
for col in cols:
    label = LabelEncoder().fit_transform(df[col])
    df.drop(col, axis=1, inplace=True)
    df[col] = label
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 5109 entries, 0 to 5109
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    5109 non-null  float64
1   hypertension            5109 non-null  int64
2   heart_disease           5109 non-null  int64
3   avg_glucose_level       5109 non-null  float64
4   bmi                     5109 non-null  float64
5   stroke                  5109 non-null  int64
6   gender                  5109 non-null  int64
7   ever_married            5109 non-null  int64
8   work_type               5109 non-null  int64
9   Residence_type          5109 non-null  int64
10  smoking_status          5109 non-null  int64
dtypes: float64(3), int64(8)
memory usage: 479.0 KB
```

Split Data

```
X = df.drop(['stroke'],axis=1)
y = df['stroke']
# split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state = 66)
```

Fit models to data

```
logit_model = LogisticRegression()  
logit_model.fit(X_train, y_train)
```

```
p = logit_model.predict(X_test)  
logit_model.score(X_test, y_test)
```

0.9500978473581213

/opt/python/envs/default/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:1043: UserWarning: STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
dt_model = DecisionTreeClassifier()  
dt_model.fit(X_train, y_train)
```

```
p = dt_model.predict(X_test)  
dt_model.score(X_test, y_test)
```

0.9060665362035225

```
rf_model = RandomForestClassifier()  
rf_model.fit(X_train, y_train)
```

```
p = rf_model.predict(X_test)  
rf_model.score(X_test, y_test)
```

0.949119373776908

```
from sklearn.metrics import make_scorer, precision_score, recall_score, accuracy_score
```

```
print(classification_report(y_test, p, target_names=['stroke', 'not']))
```

	precision	recall	f1-score	support
stroke	0.95	1.00	0.97	971
not	0.00	0.00	0.00	51
accuracy			0.95	1022
macro avg	0.48	0.50	0.49	1022
weighted avg	0.90	0.95	0.93	1022